

Technical debt, refactoring, and maintenance (1/2)

Martin Kellogg

Reading quiz: technical debt

Q1: **TRUE or FALSE:** just like financial debts, all technical debts need to eventually be paid back.

Q2: Which of the following was used as an example of a system that deferred too much maintenance and suffered for it in the article?

- A. Washington DC's Metro's signal system
- B. UC Berkeley's CalMail
- C. NASA's Mars polar orbiter
- D. AWS' us-east-1

Reading quiz: technical debt

Q1: **TRUE** or **FALSE**: just like financial debts, all technical debts need to eventually be paid back.

Q2: Which of the following was used as an example of a system that deferred too much maintenance and suffered for it in the article?

- A. Washington DC's Metro's signal system
- B. UC Berkeley's CalMail
- C. NASA's Mars polar orbiter
- D. AWS' us-east-1

Reading quiz: technical debt

Q1: **TRUE** or **FALSE**: just like financial debts, all technical debts need to eventually be paid back.

Q2: Which of the following was used as an example of a system that deferred too much maintenance and suffered for it in the article?

- A. Washington DC's Metro's signal system
- B. UC Berkeley's CalMail
- C. NASA's Mars polar orbiter
- D. AWS' us-east-1

Announcements

- Signups for preliminary demo slots (with me) are open
 - This is mandatory, but only 5/16 teams have signed up so far
 - Free reading quiz credit for all team members if you sign up for a slot before Monday morning
 - See Discord for the link
- Snafu with Wizard-of-Oz demo recordings
 - Some of you have feedback from me, some have feedback from your TA
- If you still haven't picked up your midterm, but you want to, come to office hours

Tech debt, refactoring, and maintenance (1/2)

Today's agenda:

- **Finish design pattern slides**
- Technical debt: the costs of bad design
- How to pay off technical debt: refactoring

Review: design patterns

Definition: A *software design pattern* is a general, reusable solution to a commonly occurring problem within a given context in software design.

- all patterns have **tradeoffs**. In OO languages, design patterns often trade **verbosity or efficiency** for **extensibility**

Review: design patterns

Definition: A *software design pattern* is a general, reusable solution to a commonly occurring problem within a given context in software design.

- all patterns have **tradeoffs**. In OO languages, design patterns often trade **verbosity or efficiency** for **extensibility**
- we'll consider **structural, creational and behavioral** design patterns

Review: design patterns

Definition: A *software design pattern* is a general, reusable solution to a commonly occurring problem within a given context in software design.

- all patterns have **tradeoffs**. In OO languages, design patterns often trade **verbosity or efficiency** for **extensibility**
- we'll consider **structural**, **creational** and **behavioral** design patterns
- *Structural design patterns* ease design by identifying simple ways to realize relationships among entities.
 - e.g., the adapter pattern transforms one format to another

Review: design patterns

Definition: A *software design pattern* is a general, reusable solution to a commonly occurring problem within a given context in software design.

- all patterns have **tradeoffs**. In OO languages, design patterns often trade **verbosity or efficiency** for **extensibility**
- we'll consider **structural**, **creational** and **behavioral** design patterns
- *Structural design patterns* ease design by identifying simple ways to realize relationships among entities.
 - e.g., the adapter pattern transforms one format to another
- *Creational design patterns* control object creation so that objects are created in a manner suitable for the situation.

Creational patterns: named constructor

- In the **Named Constructor Pattern**, you declare the class's normal constructors to be private or protected and make a public static creation method.

```
class Llama {  
public:  
    static Llama* create_llama(string name) {  
        return new Llama(name);  
    }  
private: // Making ctor private  
    Llama(string name_in): name(name_in) {}  
    string name;  
};
```

Creational patterns: named constructor

- In the *Named Constructor Pattern* constructors to be private or provide creation method.

```
class Llama {  
public:  
    static Llama* create_llama(string name);  
    return new Llama(name);  
}  
  
private: // Making ctor private  
    Llama(string name_in) : name(name_in) {}  
    string name;  
};
```

Why might you do this?

- might want to change to Llama subclass later
- want to validate arguments from clients, but make construction fast internally
- etc.

Creational patterns: factories

- Suppose we need to create and use polymorphic objects **without exposing their types** to the client

Creational patterns: factories

- Suppose we need to create and use polymorphic objects **without exposing their types** to the client
 - Recall: design for maintainability and extensibility. We don't want the client to depend on (and thus "lock in") the actual subtypes.

Creational patterns: factories

- Suppose we need to create and use polymorphic objects **without exposing their types** to the client
 - Recall: design for maintainability and extensibility. We don't want the client to depend on (and thus "lock in") the actual subtypes.
- The typical solution is to write a function that creates objects of the type we want but returns that object so that it appears to be ("cast to") a member of the base class

Creational patterns: factories

- Suppose we need to create and use polymorphic objects **without exposing their types** to the client
 - Recall: design for maintainability and extensibility. We don't want the client to depend on (and thus "lock in") the actual subtypes.
- The typical solution is to write a function that creates objects of the type we want but returns that object so that it appears to be ("cast to") a member of the base class
 - this is a specific variant of the named constructor pattern

Creational patterns: factories

- The *factory method pattern* (or just *factory pattern*) is a creational design pattern that uses factory methods to create objects without having the return type reveal the exact subclass created.

Creational patterns: factories

- The *factory method pattern* (or just *factory pattern*) is a creational design pattern that uses factory methods to create objects without having the return type reveal the exact subclass created.

```
Payment * payment_factory(string name, string type) {  
    if (type == "credit_card")  
        return new CreditCardPayment(name);  
    else if (type == "bitcoin")  
        return new BitcoinPayment(name);  
    ... }
```

```
Payment * webapp_session_payment =  
    payment_factory(customer_name, "credit_card");
```

Creational patterns: factories

- The **factory method pattern** (or design pattern that uses factories) without having the return type

Note how the implementation details are hidden from the client, and they can only treat the result as a **generic** payment

```
Payment * payment_factory(string name, string type) {  
    if (type == "credit_card")  
        return new CreditCardPayment(name);  
    else if (type == "bitcoin")  
        return new BitcoinPayment(name);  
    ... }
```

```
Payment * webapp_session_payment =  
payment_factory(customer_name, "credit_card");
```

Creational patterns: factories

- You may also encounter implementations in which special methods create the right type:

Creational patterns: factories

- You may also encounter implementations in which special methods create the right type:

```
class PaymentFactory {  
public:  
    static Payment* make_credit_payment(string name) {  
        return new CreditCardPayment(name);  
    }  
    static Payment* make_bc_payment(string name) {  
        return new BitcoinPayment(name);  
    } };  
Payment * webapp_session_payment =  
PaymentFactory::make_credit_payment(customer_name);
```

Creational patterns: example

- Suppose we're implementing a computer game with a **polymorphic Enemy class hierarchy**, and we want to spawn **different versions** of enemies based on the difficulty level.

Creational patterns: example

- Suppose we're implementing a computer game with a **polymorphic Enemy class hierarchy**, and we want to spawn **different versions** of enemies based on the difficulty level.

- e.g., normal difficulty = regular Goomba



- hard difficulty = spiked Goomba



Creational patterns: example: anti-patterns

- An *anti-pattern* is a common response to a recurring problem that is usually ineffective and risks being counterproductive.

Creational patterns: example: anti-patterns

- An *anti-pattern* is a common response to a recurring problem that is usually ineffective and risks being counterproductive.
- A bad solution (i.e., anti-pattern) would be to check the difficulty at each of the many places in the code related to spawning enemies:

```
Enemy* goomba = nullptr;  
if (difficulty == "normal")  
    goomba = new Goomba();  
else if (difficulty == "hard")  
    goomba = new SpikedGoomba();
```

Creational patterns: example: anti-patterns

- An *anti-pattern* is a common response to a recurring problem that is usually ineffective and risks being counterproductive.
- A bad solution (i.e., anti-pattern) would be to check the difficulty at each of the many places in the code related to spawning enemies:

```
Enemy* goomba = nullptr;  
if (difficulty == "normal")  
    goomba = new Goomba();  
else if (difficulty == "hard")  
    goomba = new SpikedGoomba();
```

Why is this bad?

Creational patterns: example: anti-patterns

- An *anti-pattern* is a common response to a recurring problem that is usually ineffective and risks being counterproductive.
- A bad solution (i.e., anti-pattern) would be to check the difficulty at each of the many places in the code related to spawning enemies:

```
Enemy* goomba = nullptr;  
if (difficulty == "normal")  
    goomba = new Goomba();  
else if (difficulty == "hard")  
    goomba = new SpikedGoomba();
```

Why is this bad?

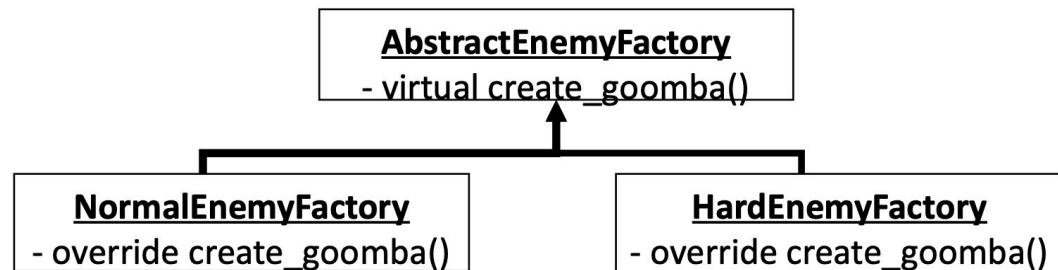
- code duplication
- consider how you'd add a new difficulty level...

Creational patterns: abstract factories

- The *abstract factory pattern* encapsulates a group of factories that have a common theme without specifying their concrete classes.

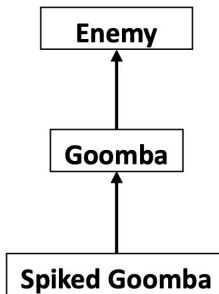
Creational patterns: abstract factories

- The *abstract factory pattern* encapsulates a group of factories that have a common theme without specifying their concrete classes.

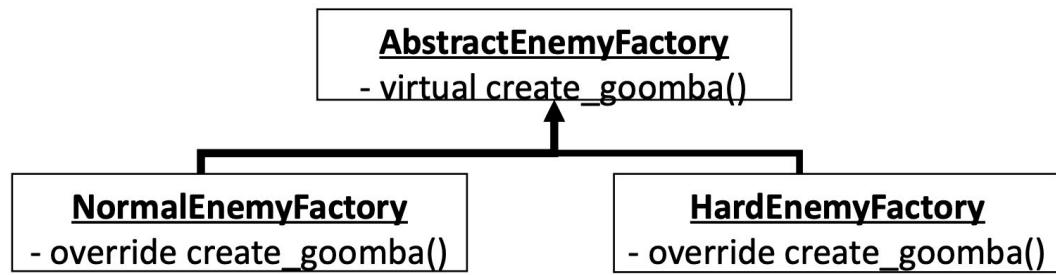


Creational patterns: abstract factories

- The *abstract factory pattern* encapsulates a group of factories that have a common theme without specifying their concrete classes.



```
// Only have to do this once!
AbstractEnemyFactory* factory = nullptr;
if (difficulty == "normal")
    factory = new NormalEnemyFactory();
else if (difficulty == "hard")
    factory = new HardEnemyFactory();
Enemy* goomba = factory->create_goomba();
```



Scenario: global application state

Scenario: global application state

- Suppose we have some application **state that needs to be globally accessible**. However, we need to control how that data is accessed and updated.

Scenario: global application state

- Suppose we have some application **state that needs to be globally accessible**. However, we need to control how that data is accessed and updated.
- The anti-pattern (**bad**) solution is to have an **unprotected global variable** (e.g., a public static field).

Scenario: global application state

- Suppose we have some application **state that needs to be globally accessible**. However, we need to control how that data is accessed and updated.
- The anti-pattern (**bad**) solution is to have an **unprotected global variable** (e.g., a public static field).
 - fails to control access or updates!

Scenario: global application state

- Suppose we have some application **state that needs to be globally accessible**. However, we need to control how that data is accessed and updated.
- The anti-pattern (**bad**) solution is to have an **unprotected global variable** (e.g., a public static field).
 - fails to control access or updates!
- A “less bad” solution is to put all of the state in one class and have a **global instance** of that class.

Scenario: global application state

- Global variables are usually a **poor design choice**. However:

Scenario: global application state

- Global variables are usually a **poor design choice**. However:
 - If you **must** access some state everywhere, passing it as a parameter to every function clutters the code (readability vs. ...)

Scenario: global application state

- Global variables are usually a **poor design choice**. However:
 - If you **must** access some state everywhere, passing it as a parameter to every function clutters the code (readability vs. ...)
 - This is not an argument for using global variables to avoid passing a few parameters.

Scenario: global application state

- Global variables are usually a **poor design choice**. However:
 - If you **must** access some state everywhere, passing it as a parameter to every function clutters the code (readability vs. ...)
 - This is not an argument for using global variables to avoid passing a few parameters.
 - Or if you need to access state stored outside your program (e.g., database, web API)

Scenario: global application state

- Global variables are usually a **poor design choice**. However:
 - If you **must** access some state everywhere, passing it as a parameter to every function clutters the code (readability vs. ...)
 - This is not an argument for using global variables to avoid passing a few parameters.
 - Or if you need to access state stored outside your program (e.g., database, web API)
 - Then global variables **may** be acceptable

Singleton design pattern

- The **singleton pattern** restricts the instantiation of a class to **exactly one** logical instance. It ensures that a class has only one logical instance at runtime and provides a global point of access to it.

Singleton

public:

- static ***get_instance()*** *// named ctor*

private:

- static ***instance*** *// the one instance*
- ***Singleton()*** *// ctor*

Singleton design pattern: example

```
class Singleton {  
    // public way to get "the one logical instance"  
    public static Singleton get_instance() {  
        if (Singleton.instance == null) Singleton.instance = new Singleton();  
        return Singleton.instance;  
    }  
    private static Singleton instance = null;  
    private Singleton() { // only runs once  
        billing_database = 0;  
        System.out.println("Singleton DB created");  
    }  
    // Our global state  
    private int billing_database;  
    public int get_billing_count() { return billing_database; }  
    public void increment_billing_count() { billing_database += 1; }  
}
```

Singleton design pattern: example

```
class Singleton {  
    // public way to get "the one logical instance"  
    public static Singleton get_instance() {  
        if (Singleton.instance == null) Singleton.instance = new Singleton();  
        return Singleton.instance;  
    }  
  
    private static Singleton instance = null;  
    private Singleton() { // only runs once  
        billing_database = 0;  
        System.out.println("Singleton DB created");  
    }  
    // Our global state  
    private int billing_database;  
    public int get_billing_count() { return billing_database; }  
    public void increment_billing_count() { billing_database += 1; }  
}
```

**lazy initializaton
of single object**

Singleton design pattern: example

```
class Singleton {  
    // public way to get "the one logical instance"  
    public static Singleton get_instance() {  
        if (Singleton.instance == null) Singleton.instance = new Singleton();  
        return Singleton.instance;  
    }  
    private static Singleton instance = null;  
    private Singleton() { // only runs once  
        billing_database = 0;  
        System.out.println("Singleton DB created");  
    }  
    // Our global state  
    private int billing_database;  
    public int get_billing_count() { return billing_database; }  
    public void increment_billing_count() { billing_database += 1; }  
}
```



this constructor
can't be called any
other way

Singleton design pattern: example

```
class Singleton {  
    // public way to get "the one logical instance"  
    public static Singleton get_instance() {  
        if (Singleton.instance == null) Singleton.instance = new Singleton();  
        return Singleton.instance;  
    }  
    private static Singleton instance = null;  
    private Singleton() { // only runs once  
        billing_database = 0;  
        System.out.println("Singleton DB created");  
    }  
    // Our global state  
    private int billing_database;  
    public int get_billing_count() { return billing_database; }  
    public void increment_billing_count() { billing_database += 1; }  
}
```

all clients share
this global state

Singleton design pattern: example

What is the output of this code?

```
class Main {  
    public static void main(String[] args) {  
        int bills = Singleton.get_instance().get_billing_count();  
        System.out.println(bills);  
  
        Singleton.get_instance().increment_billing_count();  
        bills = Singleton.get_instance().get_billing_count();  
        System.out.println(bills);  
    }  
}
```

Singleton

public:

- static **get_instance()** // *named ctor*
 - get_billing_count()
- increment_billing_count() // *adds 1*

private:

- static **instance** // *the one instance*
 - Singleton() // *ctor, prints message*
 - billing_database

Singleton design pattern: example

What is the output of this code?

```
class Main {  
    public static void main(String[] args) {  
        int bills = Singleton.get_instance().get_billing_count();  
        System.out.println(bills);  
  
        Singleton.get_instance().increment_billing_count();  
        bills = Singleton.get_instance().get_billing_count();  
        System.out.println(bills);  
    }  
}
```

Singleton

public:

- static **get_instance()** // *named ctor*
 - get_billing_count()
- increment_billing_count() // *adds 1*

private:

- static **instance** // *the one instance*
 - Singleton() // *ctor, prints message*
 - billing_database

Output:

```
Singleton DB created  
0  
1
```

Singleton design pattern: get_instance()

- Could we avoid typing Single.get_instance() so many times by doing this at all of the points in our program that use the singleton?

```
Singleton s = Singleton.get_instance();
System.out.println(s.get_billing_count());
... // later
System.out.println(s.get_billing_count());
```

Singleton design pattern: get_instance()

- Could we avoid typing Single.get_instance() so many times by doing this at all of the points in our program that use the singleton?

```
Singleton s = Singleton.get_instance();  
System.out.println(s.get_billing_count());  
... // later  
System.out.println(s.get_billing_count());
```

- Is this a good idea or not?

Singleton design pattern: get_instance()

- Could we avoid typing Single.get_instance() so many times by doing this at all of the points in our program that use the singleton?

```
Singleton s = Singleton.get_i  
System.out.println(s.get_bill  
... // later  
System.out.println(s.get_bill
```

- Is this a good idea or not?

This is a **bad idea**. There is **no guarantee** that get_instance() will return the same pointer (same object) every time it is called. (It may return different **concrete copies** of the **same logical item**.)

Singleton design pattern: another example

- Suppose we are implementing a computer version of the card game Euchre. In addition to a few abstract datatypes, we have a Game class that stores the state needed for a game of Euchre. When started, our application prototype plays one game of Euchre and then exits.
- Design question: **should we make Game a singleton?**

Singleton design pattern: another example

- Making Game a Singleton is **tempting**
 - There is only one Game instance in our application

Singleton design pattern: another example

- Making Game a Singleton is **tempting**
 - There is only one Game instance in our application
- However, there only **happens** to be one instance of Game. There's **no requirement** that we only have one instance.

Singleton design pattern: another example

- Making Game a Singleton is **tempting**
 - There is only one Game instance in our application
- However, there only **happens** to be one instance of Game. There's **no requirement** that we only have one instance.
- We should only use the Singleton pattern when current or future **requirements** dictate that only one instance should exist.

Singleton design pattern: another example

- Making Game a Singleton is **tempting**
 - There is only one Game instance in our application
- However, there only **happens** to be one instance of Game. There's **no requirement** that we only have one instance.
- We should only use the Singleton pattern when current or future **requirements** dictate that only one instance should exist.
 - Singleton is **not** a license to make everything global.

Behavioural Design Patterns

Behavioural Design Patterns

- *Behavioral design patterns* support common communication patterns among objects. They are concerned with algorithms and the assignment of responsibilities between objects.

Behavioural Design Patterns

- *Behavioral design patterns* support common communication patterns among objects. They are concerned with algorithms and the assignment of responsibilities between objects.
 - Commonly used to enable **limited sharing**

Behavioural Design Patterns

- *Behavioral design patterns* support common communication patterns among objects. They are concerned with algorithms and the assignment of responsibilities between objects.
 - Commonly used to enable **limited sharing**
 - e.g., same underlying algorithm, different interfaces or same interface, different underlying algorithms

Behavioural Design Patterns

- *Behavioral design patterns* support common communication patterns among objects. They are concerned with algorithms and the assignment of responsibilities between objects.
 - Commonly used to enable **limited sharing**
 - e.g., same underlying algorithm, different interfaces or same interface, different underlying algorithms
 - Examples: strategy pattern, template method pattern, iterator pattern, observer pattern, etc.

Iterator Pattern

- The *iterator pattern* is a common behavioral design pattern. It provides a uniform interface for traversing containers regardless of how they are implemented.

Iterator Pattern

- The *iterator pattern* is a common behavioral design pattern. It provides a uniform interface for traversing containers regardless of how they are implemented.
 - e.g., Java's List interface doesn't care whether it's backed by an array or a linked list

Iterator Pattern

- The ***iterator pattern*** is a common behavioral design pattern. It provides a uniform interface for traversing containers regardless of how they are implemented.
 - e.g., Java's List interface doesn't care whether it's backed by an array or a linked list
- Similar patterns exist for other kinds of data structures
 - e.g., ***visitor pattern*** for tree-like structures

Strategy Design Pattern

Strategy Design Pattern

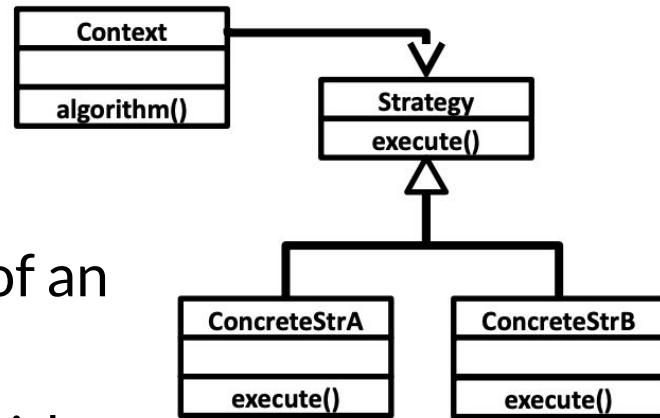
- Problem: Clients need different **variants** of an algorithm

Strategy Design Pattern

- Problem: Clients need different **variants** of an algorithm
- Solution: Create an **interface** for the algorithm, with an implementing class for each variant of the algorithm

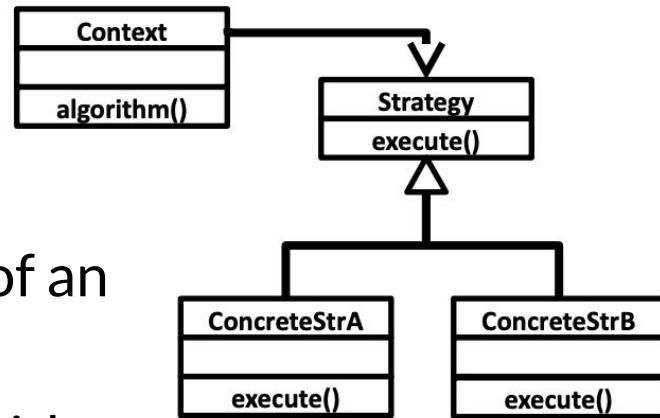
Strategy Design Pattern

- Problem: Clients need different **variants** of an algorithm
- Solution: Create an **interface** for the algorithm, with an implementing class for each variant of the algorithm



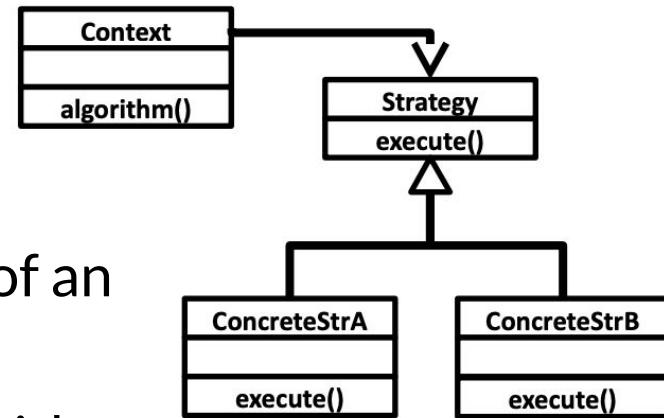
Strategy Design Pattern

- Problem: Clients need different **variants** of an algorithm
- Solution: Create an **interface** for the algorithm, with an implementing class for each variant of the algorithm
- Consequences:



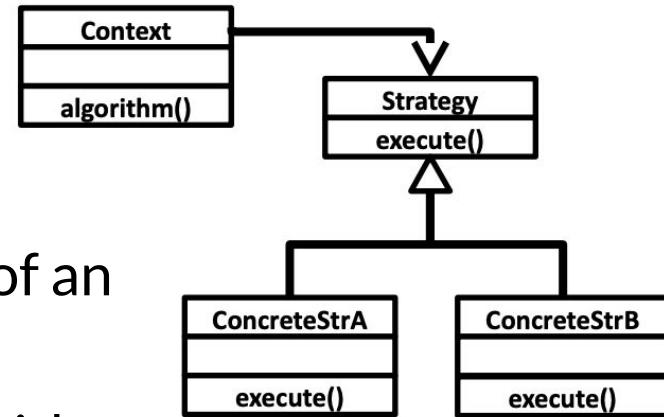
Strategy Design Pattern

- Problem: Clients need different **variants** of an algorithm
- Solution: Create an **interface** for the algorithm, with an implementing class for each variant of the algorithm
- Consequences:
 - Easily extensible for new algorithm implementations



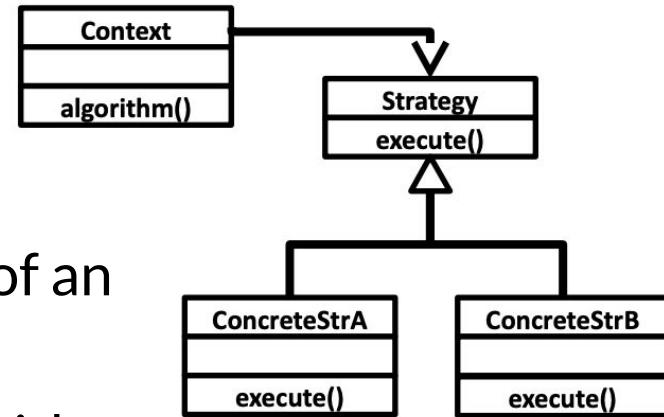
Strategy Design Pattern

- Problem: Clients need different **variants** of an algorithm
- Solution: Create an **interface** for the algorithm, with an implementing class for each variant of the algorithm
- Consequences:
 - Easily extensible for new algorithm implementations
 - Separates algorithm from client context



Strategy Design Pattern

- Problem: Clients need different **variants** of an algorithm
- Solution: Create an **interface** for the algorithm, with an implementing class for each variant of the algorithm
- Consequences:
 - Easily extensible for new algorithm implementations
 - Separates algorithm from client context
 - Introduces extra interfaces and classes: code can be harder to understand; adds overhead if the strategies are simple



Template Method Design Pattern

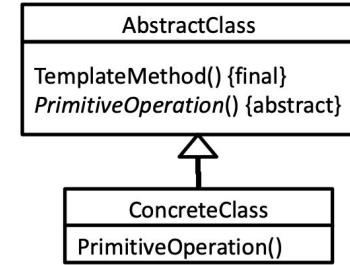
Template Method Design Pattern

- Problem: An algorithm has **customizable** and **invariant** parts

Template Method Design Pattern

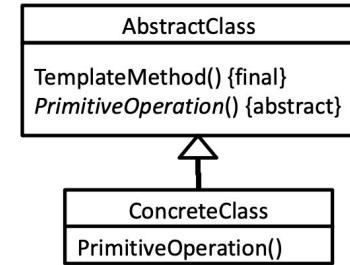
- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.

Template Method Design Pattern



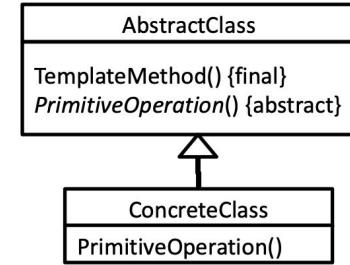
- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.

Template Method Design Pattern



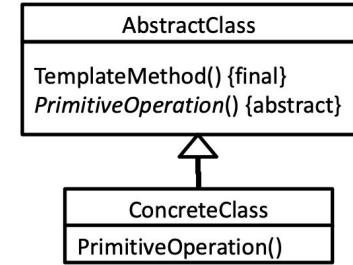
- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.
- Consequences:

Template Method Design Pattern



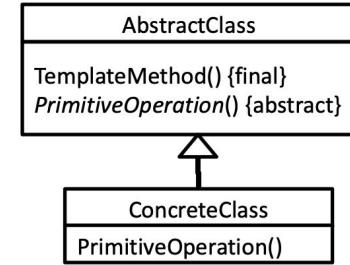
- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.
- Consequences:
 - Code reuse for the invariant parts of algorithm

Template Method Design Pattern



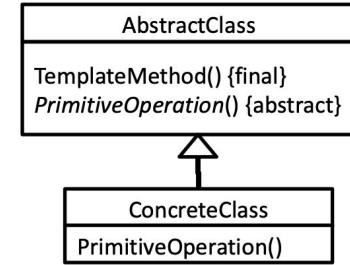
- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.
- Consequences:
 - Code reuse for the invariant parts of algorithm
 - Customization is restricted to the primitive operations

Template Method Design Pattern



- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.
- Consequences:
 - Code reuse for the invariant parts of algorithm
 - Customization is restricted to the primitive operations
 - Inverted (“Hollywood-style”) control for customization: “don’t call us, we’ll call you” (cf. comparison function in sorting)

Template Method Design Pattern



- Problem: An algorithm has **customizable** and **invariant** parts
- Solution: Implement the invariant parts of the algorithm in an **abstract** class, with abstract primitive operations representing the customizable parts of the algorithm. Subclasses customize the primitive operations.
- Consequences:
 - Code reuse for the invariant parts of algorithm
 - Customization is restricted to the primitive operations
 - Inverted (“Hollywood-style”) control for customization: “don’t call us, we’ll call you” (cf. comparison function in sorting)
 - Invariant parts of the algorithm are not changed by subclasses

Template vs. Strategy Design Pattern

Template vs. Strategy Design Pattern

- Both support variation in a larger context

Template vs. Strategy Design Pattern

- Both support variation in a larger context
- **Template method** uses inheritance + an overridable method

Template vs. Strategy Design Pattern

- Both support variation in a larger context
- **Template method** uses inheritance + an overridable method
- **Strategy** uses an interface and polymorphism (via composition)
 - Strategy objects are reusable across multiple classes
 - Multiple strategy objects are possible per class

Scenario: binge-watching

- Suppose we're implementing a video streaming website in which users can “binge-watch” (or “lock on”) to one channel. The user will then see that channel's videos in sequence. When the last such video is watched, the user should stop binge-watching that channel.

Scenario: binge-watching

- Idea: when the last video is watched, call `release_binge_watch()` on the user.

Scenario: binge-watching

- Idea: when the last video is watched, call `release_binge_watch()` on the user.

```
class User {  
    public void release_binge_watch(Channel c) {  
        if (c == binge_channel) {  
            binge_channel = null;  
        }  
    }  
    private Channel binge_channel;  
}
```

Scenario: binge-watching

- Idea: when the last video is watched, call `release_binge_watch()` on the user.

```
class User {  
    public void release_binge_watch(Channel c) {  
        if (c == binge_channel) {  
            binge_channel = null;  
        }  
    }  
    private Channel binge_channel;  
}  
  
class Channel {  
    // Called when the last video is shown  
    public void on_last_video_shown() {  
        // Global accessor for the user  
        get_user().release_binge_watch(this);  
    }  
}
```

Scenario: binge-watching

- Idea: when the last video is watched, call `release_binge_watch()` on the user.

```
class User {  
    public void release_binge_watch(Channel c) {  
        if (c == binge_channel) {  
            binge_channel = null;  
        }  
    }  
    private Channel binge_channel;  
}  
  
class Channel {  
    // Called when the last video is shown  
    public void on_last_video_shown() {  
        // Global accessor for the user  
        get_user().release_binge_watch(this);  
    }  
}
```

- What are some problems with this approach?

Scenario: binge-watching: anti-patterns

- With this design, User and Channel are **tightly coupled**
 - Changing one likely requires a change to the other

Scenario: binge-watching: anti-patterns

- With this design, User and Channel are **tightly coupled**
 - Changing one likely requires a change to the other
- The design does not support multiple users

Scenario: binge-watching: anti-patterns

- With this design, User and Channel are **tightly coupled**
 - Changing one likely requires a change to the other
- The design does not support multiple users
- What if we later want to update a user's "recommendation queue" when they finish binge-watching a channel?

Scenario: binge-watching: anti-patterns

- With this design, User and Channel are **tightly coupled**
 - Changing one likely requires a change to the other
- The design does not support multiple users
- What if we later want to update a user's "recommendation queue" when they finish binge-watching a channel?
- Whenever requirements change and we want to do something else when a video finishes (e.g., update advertising) we **must update the Channel class** and couple it to the new feature

Scenario: binge-watching: anti-patterns

- With this design, User and Channel are **tightly coupled**
 - Changing one likely requires a change to the other
- The design does not support “watch later”
- What if we later want to “recommendation queue” when they finish ~~binge-watching a channel~~:
- Whenever requirements change and we want to do something else when a video finishes (e.g., update advertising) we **must update the Channel class** and couple it to the new feature

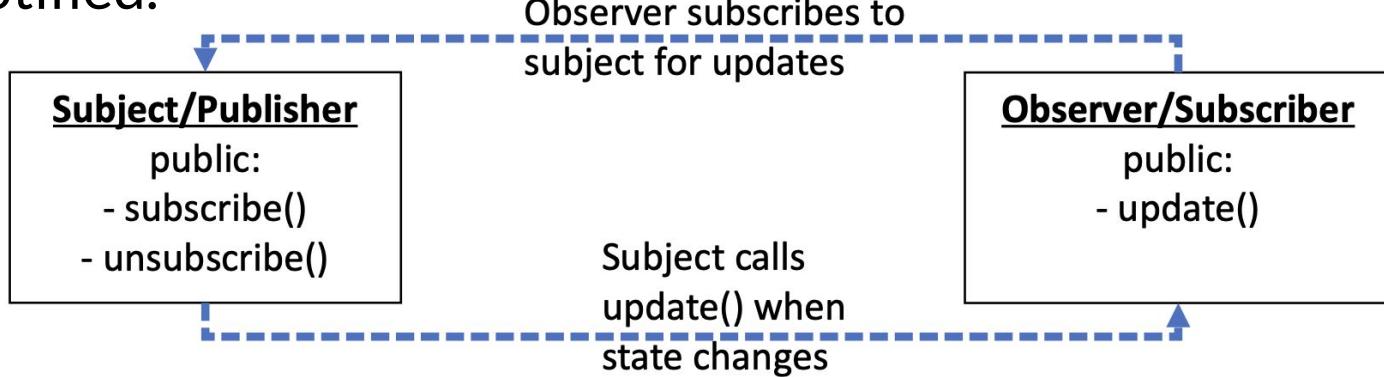
What can we do instead?

Observer Pattern

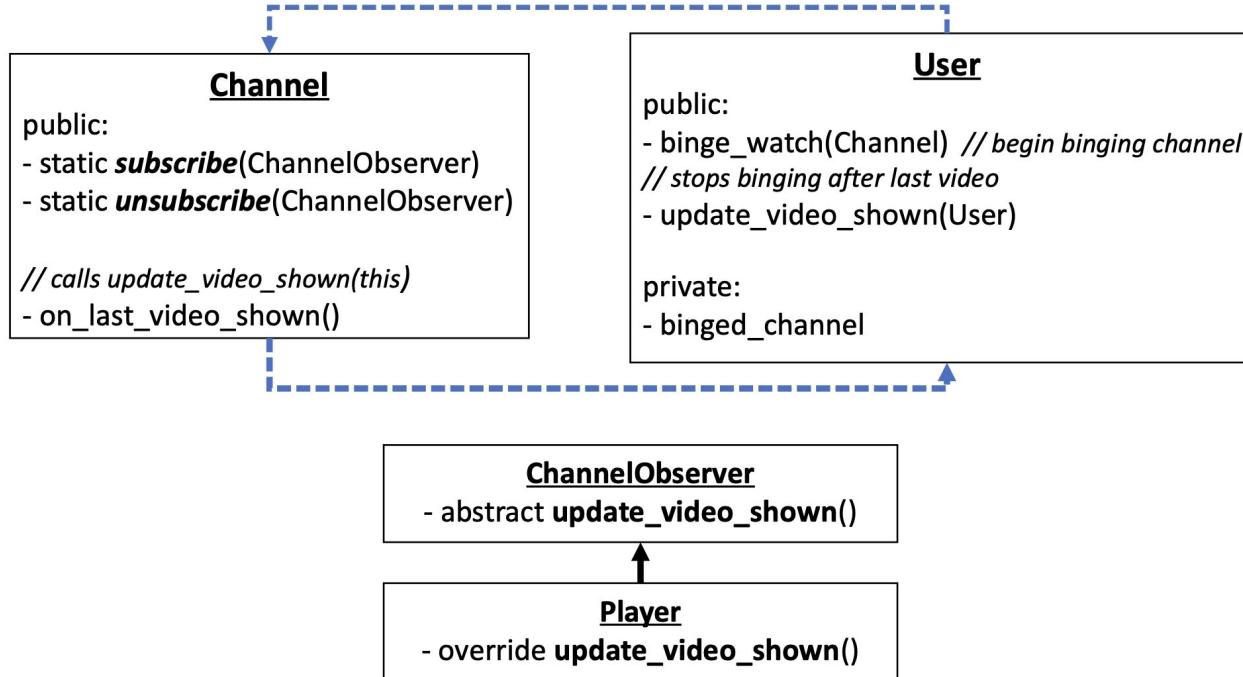
- The *observer pattern* (also called “**publish-subscribe**”) allows dependent objects to be notified automatically when the state of a subject changes. It defines a one-to-many dependency between objects so that when one object changes state, all of its dependents are notified.

Observer Pattern

- The **observer pattern** (also called “**publish-subscribe**”) allows dependent objects to be notified automatically when the state of a subject changes. It defines a one-to-many dependency between objects so that when one object changes state, all of its dependents are notified.



Observer Pattern: bing-watch scenario



Observer Pattern: bing-watch scenario

```
class Channel {  
    public static void subscribe(ChannelObserver obs) {  
        subscribers.Add(obs);  
    }  
    public static void unsubscribe(ChannelObserver obs) {  
        subscribers.Remove(obs);  
    }  
    public void on_last_video_shown() {  
        foreach (ChannelObserver obs in subscribers) {  
            observer.update_video_shown(this);  
        }  
    }  
    private static List<ChannelObserver> subscribers =  
        new List<ChannelObserver>();  
}
```

The diagram illustrates the interaction between the **User** and **Channel** classes. A dashed blue arrow points from the **User** class to the **Channel** class, indicating the direction of the communication.

Inside the **User** class, a sequence of three messages is shown:

- on_start(Channel)** // begin binging channel
- on_end(Channel)** // after last video
- on_update_video_shown(User)**

Inside the **Channel** class, a sequence of three messages is shown:

- on_subscribe(ChannelObserver)**
- on_unsubscribe(ChannelObserver)**
- on_last_video_shown()**

Observer Pattern: bing-watch scenario

```
class Channel {  
    public static void subscribe(ChannelObserver obs) {  
        subscribers.Add(obs);  
    }  
    public static void unsubscribe(ChannelObserver obs) {  
        subscribers.Remove(obs);  
    }  
    public void on_last_video_shown() {  
        foreach (ChannelObserver obs in subscribers) {  
            observer.update_video_shown(this);  
        }  
    }  
    private static List<ChannelObserver> subscribers =  
        new List<ChannelObserver>();  
}
```

```
interface ChannelObserver {  
    void update_video_shown(Channel channel);  
}
```

```
on(Channel) //begin binging channel  
y after last video  
eo_shown(User)
```

```
channel
```

```
on()
```

```
on()
```

Observer Pattern: bing-watch scenario

```
class Channel {  
    public static void subscribe(ChannelObserver obs) {  
        subscribers.Add(obs);  
    }  
    public static void unsubscribe(ChannelObserver obs) {  
        subscribers.Remove(obs);  
    }  
    public void on_last_video_shown() {  
        foreach (ChannelObserver obs in subscribers) {  
            observer.update_video_shown(this);  
        }  
    }  
    private static List<ChannelObserver> subscribers =  
        new List<ChannelObserver>();  
}
```

```
interface ChannelObserver {  
    void update_video_shown(Channel channel);  
}
```

```
(Channel) //begin binging channel  
g after last video  
eo_shown(User)  
nnel
```

```
class User: ChannelObserver {  
    public void update_video_shown(Channel c) {  
        if (c == binged_channel)  
            binged_channel = null;  
    }  
    public void binge_watch(Channel c) {  
        binged_channel = c;  
    }  
    private Channel binged_channel;  
}
```

Observer Pattern: update functions

- Having multiple “update_” functions, one for each type of state change, keeps messages **granular**

Observer Pattern: update functions

- Having multiple “update_” functions, one for each type of state change, keeps messages **granular**
 - Observers that do not care about a particular type of update can ignore it (via an empty implementation of the update function)

Observer Pattern: update functions

- Having multiple “update_” functions, one for each type of state change, keeps messages **granular**
 - Observers that do not care about a particular type of update can ignore it (via an empty implementation of the update function)
- Generally it is better to pass the newly-updated data as a parameter to the update function (**push**) as opposed to making observers fetch it each time (**pull**)

Design patterns: takeaways

- Thinking about design before you start coding is usually worthwhile for large projects
 - Design around the most expensive parts of the software engineering process (usually maintenance!)
- Design patterns are reusable solutions to common problems
- Be familiar with them enough to recognize when they're being used
 - and to know when to use them yourself
 - you can look up details of a pattern if you remember its name!
- Be mindful of and avoid common anti-patterns

Tech debt, refactoring, and maintenance (1/2)

Today's agenda:

- Finish design pattern slides
- **Technical debt: the costs of bad design**
- How to pay off technical debt: refactoring

Technical debt

Technical debt

Definition: a *technical debt* is a sub-optimal design decision taken intentionally in order to gain some immediate benefit

Technical debt

Definition: a *technical debt* is a sub-optimal design decision taken intentionally in order to gain some immediate benefit

- analogy to **financial debts**:

Technical debt

Definition: a *technical debt* is a sub-optimal design decision taken intentionally in order to gain some immediate benefit

- analogy to **financial debts**:
 - you gain some immediate benefit
 - in a financial debt, you gain a large sum of money
 - in a technical debt, you gain implementation speed, etc.

Technical debt

Definition: a *technical debt* is a sub-optimal design decision taken intentionally in order to gain some immediate benefit

- analogy to **financial debts**:
 - you gain some immediate benefit
 - in a financial debt, you gain a large sum of money
 - in a technical debt, you gain implementation speed, etc.
 - you pay for it over time
 - in a financial debt, you pay interest
 - in a technical debt, your maintenance costs increase

Technical debt: benefits

- Why might you **intentionally** make a sub-optimal design decision?

Technical debt: benefits

- Why might you **intentionally** make a sub-optimal design decision?
 - Cost
 - either in dev time or because the code isn't done yet
 - Need to meet a deadline
 - Avoid premature optimization
 - Code reuse
 - Principle of least surprise
 - Organizational requirements/politics
 - etc.

Technical debt: paying interest

- Unlike a financial debt, a technical debt doesn't have a **creditor**

Technical debt: paying interest

- Unlike a financial debt, a technical debt doesn't have a **creditor**
 - Conceptually, when you take on technical debt you are borrowing from **future maintainers** of the system

Technical debt: paying interest

- Unlike a financial debt, a technical debt doesn't have a **creditor**
 - Conceptually, when you take on technical debt you are borrowing from **future maintainers** of the system
- Recall our goals in good design:

Technical debt: paying interest

- Unlike a financial debt, a technical debt doesn't have a **creditor**
 - Conceptually, when you take on technical debt you are borrowing from **future maintainers** of the system
- Recall our goals in good design:
 - design for **change and reuse**
 - make the system easy to extend, modify, etc.

Technical debt: paying interest

- Unlike a financial debt, a technical debt doesn't have a **creditor**
 - Conceptually, when you take on technical debt you are borrowing from **future maintainers** of the system
- Recall our goals in good design:
 - design for **change and reuse**
 - make the system easy to extend, modify, etc.
- **Implication:** a system with technical debt is **harder** to change and reuse

Technical debt: benefits and costs

Examples of debt:

Examples of costs:

Technical debt: benefits and costs

Examples of debt:

- code smells

Examples of costs:

- “smelly” code is less flexible

Technical debt: benefits and costs

Examples of debt:

- code smells
- missing tests

Examples of costs:

- “smelly” code is less flexible
- tests don’t catch breaking change, causing outages

Technical debt: benefits and costs

Examples of debt:

- code smells
- missing tests
- missing documentation

Examples of costs:

- “smelly” code is less flexible
- tests don’t catch breaking change, causing outages
- need to spend time to figure out how system works

Technical debt: benefits and costs

Examples of debt:

- code smells
- missing tests
- missing documentation
- dependency on old versions of third-party systems

Examples of costs:

- “smelly” code is less flexible
- tests don’t catch breaking change, causing outages
- need to spend time to figure out how system works
- may need to take over maintenance of old system

Technical debt: when is it worth it?

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?
 - e.g., safety, performance, scalability, etc.

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?
 - e.g., safety, performance, scalability, etc.
 - And how do our architectural decisions reflect those attributes?

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?
 - e.g., safety, performance, scalability, etc.
 - And how do our architectural decisions reflect those attributes?
 - i.e., will we be able to reach our goals using this design?

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?
 - e.g., safety, performance, scalability, etc.
 - And how do our architectural decisions reflect those attributes?
 - i.e., will we be able to reach our goals using this design?
- The choice to take on technical debt is always a **tradeoff**:

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?
 - e.g., safety, performance, scalability, etc.
 - And how do our architectural decisions reflect those attributes?
 - i.e., will we be able to reach our goals using this design?
- The choice to take on technical debt is always a **tradeoff**:
 - give up some flexibility later, gain something now

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality attributes** that our software needs to ultimately satisfy?
 - e.g., safety, performance, scalability, etc.
 - And how do our architectural decisions reflect those attributes?
 - i.e., will we be able to reach our goals using this design?
- The choice to take on technical debt is always a **tradeoff**:
 - give up some flexibility later, gain something now
 - whether this is worthwhile varies **case by case**

Technical debt: when is it worth it?

- Key consideration:
 - What are the **quality** requirements we ultimately satisfy?
 - e.g., safety, performance, etc.
 - And how do our architectural choices **contribute** to those requirements?
 - i.e., will we be able to reach our goals using this design?
- The choice to take on technical debt is always a **tradeoff**:
 - give up some flexibility later, gain something now
 - whether this is worthwhile varies **case by case**

Whether to take on technical debt is often one of the **most consequential** choices you get to make as an engineer. **Take it seriously!**

Technical debt: when is it worth it?

- You should also consider **risk** when taking on technical debt

Technical debt: when is it worth it?

- You should also consider **risk** when taking on technical debt
 - i.e., ask yourself “what is the **worst thing** that could happen in the future if I take this shortcut today”?

Technical debt: when is it worth it?

- You should also consider **risk** when taking on technical debt
 - i.e., ask yourself “what is the **worst thing** that could happen in the future if I take this shortcut today”?
 - risk should preclude you from taking on certain kind of debts
 - e.g., never use laughably-bad security or break laws, even if you don’t plan to deploy this prototype

Technical debt: when is it worth it?

- You should also consider **risk** when taking on technical debt
 - i.e., ask yourself “what is the **worst thing** that could happen in the future if I take this shortcut today”?
 - risk should preclude you from taking on certain kind of debts
 - e.g., never use laughably-bad security or break laws, even if you don’t plan to deploy this prototype
- Best practice (especially for relatively risky debts): **write everything down!**
 - that way, you know what you need to fix before releasing

Technical debt: Y2k example

- History quiz: what was the “**Y2k bug**”?

Technical debt: Y2k example

- History quiz: what was the “**Y2k bug**”?
 - Answer: many early programs stored the year using **two digits**
 - assumption: current year = “19” + those two digits

Technical debt: Y2k example

- History quiz: what was the “**Y2k bug**”?
 - Answer: many early programs stored the year using **two digits**
 - assumption: current year = “19” + those two digits
- This is an example of technical debt:

Technical debt: Y2k example

- History quiz: what was the “**Y2k bug**”?
 - Answer: many early programs stored the year using **two digits**
 - assumption: current year = “19” + those two digits
- This is an example of technical debt:
 - **immediate benefit**: saves hard disk space (expensive in 1980)

Technical debt: Y2k example

- History quiz: what was the “**Y2k bug**”?
 - Answer: many early programs stored the year using **two digits**
 - assumption: current year = “19” + those two digits
- This is an example of technical debt:
 - **immediate benefit**: saves hard disk space (expensive in 1980)
 - **long-term cost**: if the program is still being used in 2000, need to fix it!
 - “I just never imagined anyone would be using these systems 10 years later, let alone 20.”

Technical debt: not always strictly technical

- You can also view **other serious risks** to the system's continued maintenance as forms of technical debt

Technical debt: not always strictly technical

- You can also view **other serious risks** to the system's continued maintenance as forms of technical debt
 - e.g., if your **bus factor** (= “number of people who need to get hit by a bus before no one understands the system”) is low and parts of the system are undocumented...

Technical debt: not always strictly technical

- You can also view **other serious risks** to the system's continued maintenance as forms of technical debt
 - e.g., if your **bus factor** (= “number of people who need to get hit by a bus before no one understands the system”) is low and parts of the system are undocumented...
 - the amount of technical debt you have is higher than if your bus factor was very high

Technical debt: not always strictly technical

- You can also view **other serious risks** to the system's continued maintenance as forms of technical debt
 - e.g., if your **bus factor** (= “number of people who need to get hit by a bus before no one understands the system”) is low and parts of the system are undocumented...
 - the amount of technical debt you have is higher than if your bus factor was very high
- Other examples include having **high staff turnover** (which systematically lowers bus factor) or few senior engineers

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase that already exists

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase that already exists
 - we usually call such a codebase *legacy code*

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase that already exists
 - we usually call such a codebase *legacy code*
- What if this code **already** has technical debt? (Hint: it **always** does.)

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase that already exists
 - we usually call such a codebase *legacy code*
- What if this code **already** has technical debt? (Hint: it **always** does.)
 - You **must service** the debt: you must deal with the code as it is

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase that already exists
 - we usually call such a codebase *legacy code*
- What if this code **already** has technical debt? (Hint: it **always** does.)
 - You **must service** the debt: you must deal with the code as it is
 - You **do not gain** the benefit: the benefit was immediate, but you're reaching the code too late to see it

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase
 - we usually
 - What if this codebase is bad?
 - You **must** support it
 - You **do not** have time to fix it, but you're reading this

Unfortunate but common anti-pattern:

ys does as it is, but

Unfortunate but common anti-pattern:

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase
 - we usually
- What if this codebase
 - You **must** support it
 - You **do not** understand it
you're reading this

Unfortunate but common anti-pattern:

- dev 1 builds a new system, taking on a lot of technical debt

(**sys** does.)
as it is,
, but

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase
 - we usually
- What if this codebase
 - You **must** support it
 - You **do not** have time to fix it, so you're reading this

Unfortunate but common anti-pattern:

- dev 1 builds a new system, taking on a lot of technical debt
- system is successful initially, dev 1 is promoted or moves on

(**sys** does.)
as it is,
, but

Technical debt: not always your fault

- Common situation: you are now responsible for maintaining and improving a codebase
 - we usually
- What if this codebase
 - You **must** support it
 - You **do not** understand it
you're reading this slide

Unfortunate but common anti-pattern:

- dev 1 builds a new system, taking on a lot of technical debt
- system is successful initially, dev 1 is promoted or moves on
- dev 2 is now responsible for paying the debt on the system :(

ys does.)
e as it is,
, but

Technical debt: bitrot

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally
 - even if the code was initially reviewed and well-designed at the time of commit, and even if changes are reviewed, etc.

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally
 - even if the code was initially reviewed and well-designed at the time of commit, and even if changes are reviewed, etc.
 - this process is called “**bitrot**”

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally
 - even if the code was initially reviewed and well-designed at the time of commit, and even if changes are reviewed, etc.
 - this process is called “**bitrot**”
- Why does bitrot happen?

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally
 - even if the code was initially reviewed and well-designed at the time of commit, and even if changes are reviewed, etc.
 - this process is called “**bitrot**”
- Why does bitrot happen?
 - Systems evolve to meet new needs and add new features

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally
 - even if the code was initially reviewed and well-designed at the time of commit, and even if changes are reviewed, etc.
 - this process is called “**bitrot**”
- Why does bitrot happen?
 - Systems evolve to meet new needs and add new features
 - Changes happen in dependencies, languages, environment

Technical debt: bitrot

- Over time, software tends to have **increasing maintenance costs**, even if no technical debt is taken on intentionally
 - even if the code was initially reviewed and well-designed at the time of commit, and even if changes are reviewed, etc.
 - this process is called “**bitrot**”
- Why does bitrot happen?
 - Systems evolve to meet new needs and add new features
 - Changes happen in dependencies, languages, environment
 - If the code's structure does not also evolve, it will "rot"

Technical debt example: languages

- **Language choice** is a common example of a place where it might make sense to take on technical debt:

Technical debt example: languages

- **Language choice** is a common example of a place where it might make sense to take on technical debt:
 - relatively-unsafe and/or non-performant languages (e.g., Python, Ruby, JavaScript) are **easier to write** code in

Technical debt example: languages

- **Language choice** is a common example of a place where it might make sense to take on technical debt:
 - relatively-unsafe and/or non-performant languages (e.g., Python, Ruby, JavaScript) are **easier to write** code in
 - but, if you end up needing to write performance-critical or safety-critical code in them, you're going to have a bad time!

Technical debt example: languages

- **Language choice** is a common example of a place where it might make sense to take on technical debt:
 - relatively-unsafe and/or non-performant languages (e.g., Python, Ruby, JavaScript) are **easier to write** code in
 - but, if you end up needing to write performance-critical or safety-critical code in them, you're going to have a bad time!
 - on the other hand, investing in writing in a safe and performant language (e.g., Rust, Kotlin) has a **higher upfront cost**

Technical debt example: languages

- **Language choice** is a common example of a place where it might make sense to take on technical debt:
 - relatively-unsafe and/or non-performant languages (e.g., Python, Ruby, JavaScript) are **easier to write** code in
 - but, if you end up needing to write performance-critical or safety-critical code in them, you're going to have a bad time!
 - on the other hand, investing in writing in a safe and performant language (e.g., Rust, Kotlin) has a **higher upfront cost**
 - but you might save a big headache later

Technical debt example: languages

- **Language choice** is a common example of a place where it might make sense to take on technical debt:
 - relatively-unsafe and/or non-performant languages (e.g., Python, Ruby)
 - but, if you're safety-conscious, it might be better to use a language that's more safe and performant)
 - on the other hand, you might choose a less safe or less performant language (e.g., C++)
 - but you might save a big headache later

Technical debt example: Facebook + PHP

- Facebook's original site was written in PHP in 2004

Technical debt example: Facebook + PHP

- Facebook's original site was written in PHP in 2004
 - PHP is dynamically-typed and **relatively unsafe**

Technical debt example: Facebook + PHP

- Facebook's original site was written in PHP in 2004
 - PHP is dynamically-typed and **relatively unsafe**
 - this caused problems for Facebook as its codebase grew

Technical debt example: Facebook + PHP

- Facebook's original site was written in PHP in 2004
 - PHP is dynamically-typed and **relatively unsafe**
 - this caused problems for Facebook as its codebase grew
- In 2014, Facebook releases **Hack**, a new variant of PHP

Technical debt example: Facebook + PHP

- Facebook's original site was written in PHP in 2004
 - PHP is dynamically-typed and **relatively unsafe**
 - this caused problems for Facebook as its codebase grew
- In 2014, Facebook releases Hack, a new variant of PHP
 - Hack added **new safety features** (including gradual typing and type inference)

Technical debt example: Facebook + PHP

- Facebook's original site was written in PHP in 2004
 - PHP is dynamically-typed and **relatively unsafe**
 - this caused problems for Facebook as its codebase grew
- In 2014, Facebook releases Hack, a new variant of PHP
 - Hack added **new safety features** (including gradual typing and type inference)
 - “Hack enables us to dynamically convert our code one file at a time” - Facebook Technical Lead, HipHop VM (HHVM)

Technical debt example: machine learning

Technical debt example: machine learning

- Machine-learning components can encourage tech debt

Technical debt example: machine learning

- Machine-learning components can encourage tech debt
 - hard to enforce **strict abstraction boundaries**

Technical debt example: machine learning

- Machine-learning components can encourage tech debt
 - hard to enforce **strict abstraction boundaries**
 - after all, one big reason for ML is that the desired behavior **cannot be effectively implemented in software logic** without dependency on external data!

Technical debt example: machine learning

- Machine-learning components can encourage tech debt
 - hard to enforce **strict abstraction boundaries**
 - after all, one big reason for ML is that the desired behavior **cannot be effectively implemented in software logic** without dependency on external data!
- For this reason, Google engineers have called ML systems the “**high-interest credit card**” of technical debt [1]

Technical debt example: machine learning

- Machine-learning components can encourage tech debt
 - hard to enforce **strict abstraction boundaries**
 - after all, one big reason for ML is that the desired behavior **cannot be effectively implemented in software logic** without dependency on external data!
- For this reason, Google engineers have called ML systems the “**high-interest credit card**” of technical debt [1]
 - can get you a lot of value in the short term!

Technical debt example: machine learning

- Machine-learning components can encourage tech debt
 - hard to enforce **strict abstraction boundaries**
 - after all, one big reason for ML is that the desired behavior **cannot be effectively implemented in software logic** without dependency on external data!
- For this reason, Google engineers have called ML systems the “**high-interest credit card**” of technical debt [1]
 - can get you a lot of value in the short term!
 - but if you don’t pay down the debt quickly...

Topical aside: LLMs and technical debt

Topical aside: LLMs and technical debt

- It is **not yet well understood** how LLM code generators (e.g., GitHub CoPilot and friends) interact with technical debt

Topical aside: LLMs and technical debt

- It is **not yet well understood** how LLM code generators (e.g., GitHub CoPilot and friends) interact with technical debt
- However, early signs are **not promising**:

Topical aside: LLMs and technical debt

- It is **not yet well understood** how LLM code generators (e.g., GitHub CoPilot and friends) interact with technical debt
- However, early signs are **not promising**:
 - LLMs seem to be easily confused by atypical code patterns, quirks of leaky abstractions, etc. (all hallmarks of tech debt)

Topical aside: LLMs and technical debt

- It is **not yet well understood** how LLM code generators (e.g., GitHub CoPilot and friends) interact with technical debt
- However, early signs are **not promising**:
 - LLMs seem to be easily confused by atypical code patterns, quirks of leaky abstractions, etc. (all hallmarks of tech debt)
 - LLMs can **introduce technical debt** themselves
 - e.g., recent studies have shown that with an LLM assistant, devs are more likely to write insecure code [1]

[1] Do Users Write More Insecure Code with AI Assistants? Neil Perry, Megha Srivastava, Deepak Kumar, and Dan Boneh. CCS 2023.

Paying down technical debt

- It is possible to **reduce** the amount of technical debt in a codebase by improving its design

Paying down technical debt

- It is possible to **reduce** the amount of technical debt in a codebase by improving its design
 - one option: **rewriting** the whole system (but think about today's Spolsky reading!)

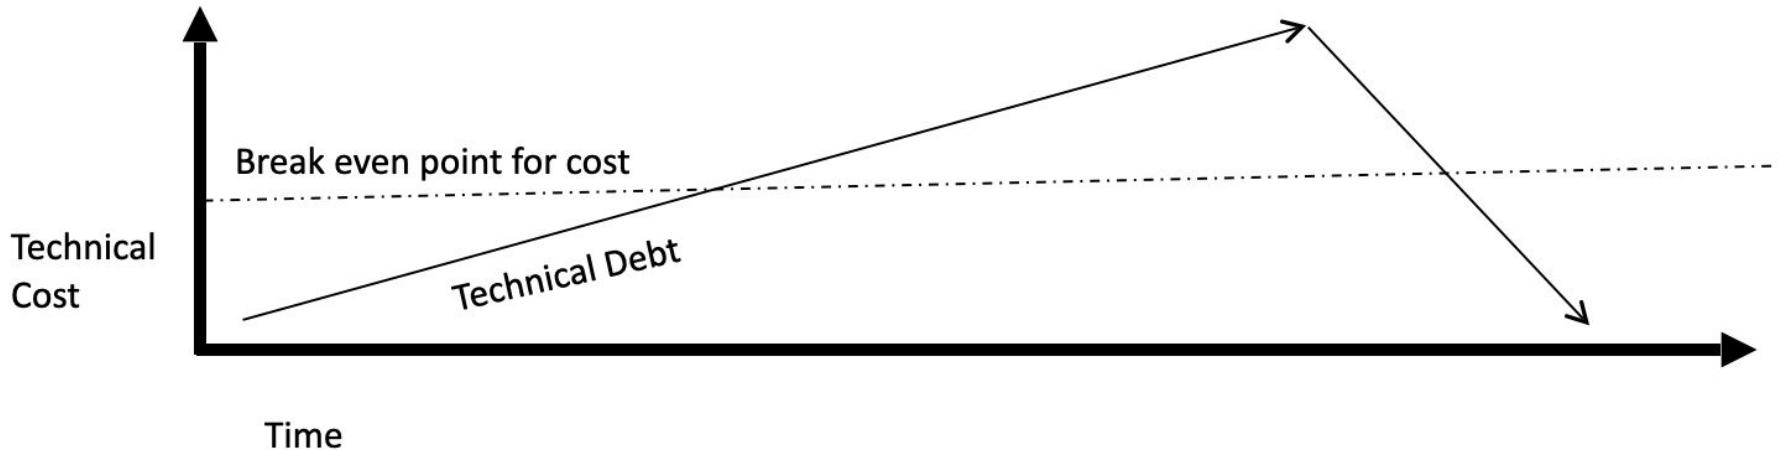
Paying down technical debt

- It is possible to **reduce** the amount of technical debt in a codebase by improving its design
 - one option: **rewriting** the whole system (but think about today's Spolsky reading!)
 - more common: **refactoring** the code

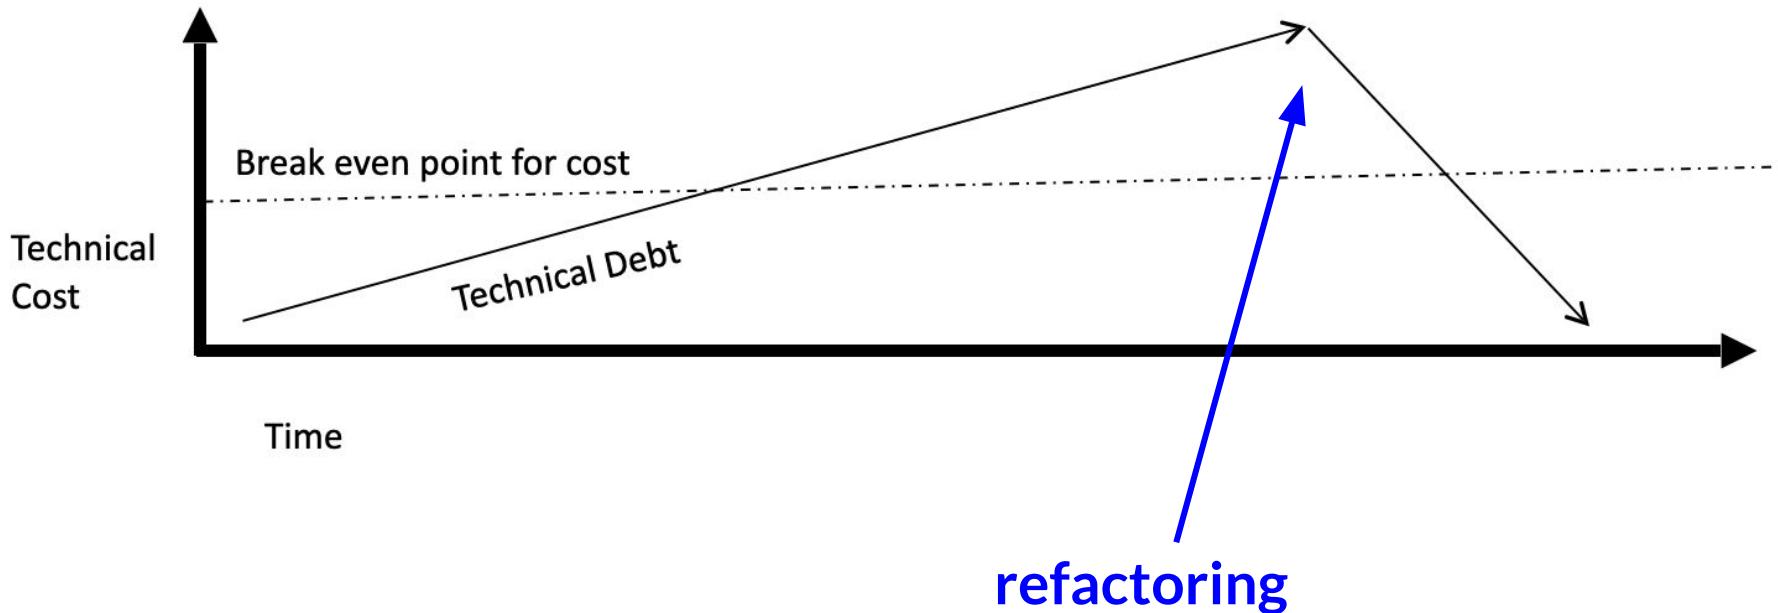
Paying down technical debt

- It is possible to **reduce** the amount of technical debt in a codebase by improving its design
 - one option: **rewriting** the whole system (but think about today's Spolsky reading!)
 - more common: **refactoring** the code
- **refactoring** is the process of applying behaviour-preserving transformations (called **refactorings**) to a program, with the goal of improving its non-functional properties (e.g., design, performance)

Paying down technical debt



Paying down technical debt



Paying down technical debt: best practices

Paying down technical debt: best practices

- Advice: set aside **specific time** to pay off technical debt

Paying down technical debt: best practices

- Advice: set aside **specific time** to pay off technical debt
 - Google has (had?) “20% time” for tasks like this

Paying down technical debt: best practices

- Advice: set aside **specific time** to pay off technical debt
 - Google has (had?) “20% time” for tasks like this
- **New projects** can take on some technical debt

Paying down technical debt: best practices

- Advice: set aside **specific time** to pay off technical debt
 - Google has (had?) “20% time” for tasks like this
- **New projects** can take on some technical debt
 - i.e., refactoring at the start of a project to make the rest of the new code easier to write

Paying down technical debt: best practices

- Advice: set aside **specific time** to pay off technical debt
 - Google has (had?) “20% time” for tasks like this
- **New projects** can take on some technical debt
 - i.e., refactoring at the start of a project to make the rest of the new code easier to write
- Have a plan: **don’t put off dealing with technical debt indefinitely**

Paying down technical debt: best practices

- Advice: set aside **specific time** to pay off technical debt
 - Google has (had?) “20% time” for tasks like this
- **New projects** can take on some technical debt
 - i.e., refactoring at the start of a project to make the rest of the new code easier to write
- Have a plan: **don’t put off dealing with technical debt indefinitely**
 - When a crisis hits, it’s too late
 - Hasty fixes to unmaintainable code likely to multiply problems!
 - Eventually, mounting technical debt can bury a team

Tech debt, refactoring, and maintenance

Agenda:

- Finish design pattern slides
- Technical debt: the costs of bad design
- **How to pay off technical debt: refactoring**

Refactoring

Definition: *refactoring* is improving a piece of software's internal structure without altering its external behavior.

Refactoring

Definition: *refactoring* is improving a piece of software's internal structure without altering its external behavior.

- Incurs a short-term time/work cost to reap **long-term benefits**

Refactoring

Definition: *refactoring* is improving a piece of software's internal structure without altering its external behavior.

- Incurs a short-term time/work cost to reap **long-term benefits**
- A long-term **investment** in the overall quality of your system.

Refactoring

Definition: *refactoring* is improving a piece of software's internal structure without altering its external behavior.

- Incurs a short-term time/work cost to reap **long-term benefits**
- A long-term **investment** in the overall quality of your system.

What refactoring is **not**:

Refactoring

Definition: *refactoring* is improving a piece of software's internal structure without altering its external behavior.

- Incurs a short-term time/work cost to reap **long-term benefits**
- A long-term **investment** in the overall quality of your system.

What refactoring is **not**:

- rewriting code
- adding features
- debugging code

Refactoring: motivation

Question: why fix a part of your system that **isn't broken?**

Refactoring: motivation

Question: why fix a part of your system that **isn't broken?**

- Each part of your system's code has three purposes:
 - to execute its functionality,
 - to allow change,
 - to communicate well to developers who read it.

Refactoring: motivation

Question: why fix a part of your system that **isn't broken?**

- Each part of your system's code has three purposes:
 - to execute its functionality,
 - to allow change,
 - to communicate well to developers who read it.
- If the code does not do one or more of these, it **is** broken.

Refactoring: motivation

Question: why fix a part of your system that **isn't broken?**

- Each part of your system's code has three purposes:
 - to execute its functionality,
 - to allow change,
 - to communicate well to developers who read it.
- If the code does not do one or more of these, it **is** broken.
- Refactoring should improve the software's design:
 - more extensible, flexible, understandable, performant, ...
 - every design improvement has costs (and risks)

Refactoring: when to refactor

Refactoring: when to refactor

Definition: a “*code smell*” is a minor design issue with a piece of code that is not a defect *per se*, but is still undesirable

Refactoring: when to refactor

Definition: a “**code smell**” is a minor design issue with a piece of code that is not a defect *per se*, but is still undesirable

- intuition: each code smell is an **irritation** on its own, but in large groups they impede maintenance

Refactoring: when to refactor

Definition: a “**code smell**” is a minor design issue with a piece of code that is not a defect *per se*, but is still undesirable

- intuition: each code smell is an **irritation** on its own, but in large groups they impede maintenance
- many code smells -> good idea to refactor

Refactoring: when to refactor

Definition: a “**code smell**” is a minor design issue with a piece of code that is not a defect *per se*, but is still undesirable

- intuition: each code smell is an **irritation** on its own, but in large groups they impede maintenance
- many code smells -> good idea to refactor
- a good refactoring often fixes more than one code smell
 - sometimes many more than one

Refactoring: when to refactor

Examples of **common code smells**:

Refactoring: when to refactor

Examples of **common code smells**:

- Duplicated code
- Poor abstraction (change one place → must change others)
- Large loop, method, class, parameter list; deeply nested loop
- Module has too little cohesion
- Modules have too much coupling
- Module has poor encapsulation
- Dead code
- Design is unnecessarily general
- Design is too specific

Refactoring: “low-level” refactoring

- “*low-level*” refactorings are small changes to the code that mitigate or remove one or more code smells. Examples:

Refactoring: “low-level” refactoring

- “*low-level*” refactorings are small changes to the code that mitigate or remove one or more code smells. Examples:
 - Renaming (methods, variables)
 - Naming (extracting) “magic” constants
 - Extracting common functionality (including duplicate code) into a module/method/etc.
 - Changing method signatures
 - Splitting one method into two or more to improve cohesion and readability (by reducing its size)

also see <https://refactoring.com/catalog/>

Refactoring: “low-level” refactoring

- modern IDEs have good support for low-level refactoring

Refactoring: “low-level” refactoring

- modern IDEs have good support for low-level refactoring
 - *IDE* = “*integrated development environment*”
 - e.g., Eclipse, VSCode, IntelliJ, etc.

Refactoring: “low-level” refactoring

- modern IDEs have good support for low-level refactoring
 - *IDE* = “*integrated development environment*”
 - e.g., Eclipse, VSCode, IntelliJ, etc.
- they automate:
 - renaming of variables, methods, classes
 - extraction of methods and constants
 - extraction of repetitive code snippets
 - changing method signatures
 - warnings about inconsistent code
 - ...

Refactoring: “low-level” refactoring

- modern IDEs have good support for low-level refactoring
 - **IDE** = “*integrated development environment*”
 - e.g., Eclipse, VSCode, IntelliJ etc
- they automate:
 - renaming of variables, methods, classes, etc.
 - extraction of methods and classes
 - extraction of repetitive code
 - changing method signatures
 - warnings about inconsistent code
 - ...

My advice/opinion: don't rely on your IDE too much. It's useful for auto-complete, simple refactoring, red squiggles, etc. But, if you let it control the build process you'll have a bad time.

Refactoring: “high-level” refactoring

- “*High-level*” refactoring might include:

Refactoring: “high-level” refactoring

- “*High-level*” refactoring might include:
 - Refactoring to design patterns
 - Changing language idioms (safety, brevity)
 - Performance optimization
 - Clarifying a statement that has evolved over time or is unclear

Refactoring: “high-level” refactoring

- “*High-level*” refactoring might include:
 - Refactoring to design patterns
 - Changing language idioms (safety, brevity)
 - Performance optimization
 - Clarifying a statement that has evolved over time or is unclear
- Compared to low-level refactoring, high-level is:

Refactoring: “high-level” refactoring

- “*High-level*” refactoring might include:
 - Refactoring to design patterns
 - Changing language idioms (safety, brevity)
 - Performance optimization
 - Clarifying a statement that has evolved over time or is unclear
- Compared to low-level refactoring, high-level is:
 - Not as well-supported by tools
 - But much **more important!**

Refactoring: how to refactor

- When you identify an area of your system that:

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features...**

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features**...

These are a good set of criteria for deciding to refactor code

- especially “needs new features”, because if you don’t refactor you’ll be **paying interest** on the tech debt!

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features...**
- What should you do?

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features...**
- What should you do?
 - Write **unit tests** that verify the code's external correctness.
(They should pass on the current, badly-designed code.)

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features...**
- What should you do?
 - Write **unit tests** that verify the code's external correctness.
(They should pass on the current, badly-designed code.)
 - **Refactor** the code. (Some unit tests may break. Fix the bugs.)

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features**...
- What should you do?
 - Write **unit tests** that verify the code's external correctness.
(They should pass on the current, badly-designed code.)
 - **Refactor** the code. (Some unit tests may break. Fix the bugs.)
 - Add any **new features**.

Refactoring: how to refactor

- When you identify an area of your system that:
 - is **poorly designed**, and
 - is **poorly tested** (even if it seems to work so far), and
 - now **needs new features**...
- What should you do?
 - Write **unit tests** that verify the code's external correctness.
(They should pass on the current, badly-designed code.)
 - **Refactor** the code. (Some unit tests may break. Fix the bugs.)
 - Add any **new features**.
 - As always, keep changes small, do code reviews, etc.

Takeaways: tech debt and refactoring

- Technical debt accrues when you take a shortcut for some immediate benefit that makes a system harder to maintain
 - tech debt is inevitable in large systems
 - but you should be thoughtful about when/how you take it on!
- When and how you take on technical debt is one of the biggest judgment calls that you will make as a low-level engineer
- Refactoring is the process of improving a codebase's non-functional properties while maintaining its behavior
 - refactoring is a useful way to reduce tech debt
 - you often want to pair refactoring with adding new features